

Python for Machine Learning

A Complete Learning Journey

Mayur Patil

August 23, 2026

Contents

I	Python Fundamentals	5
1	Python Data Types	7
1.1	Why This Matters	7
1.2	Explanation	7
1.2.1	1. Numeric Types	7
1.2.2	2. Text and Binary Types	7
1.2.3	3. Boolean and None	7
1.2.4	4. Collection Types (Brief intro)	7
1.2.5	5. <code>type()</code> and <code>isinstance()</code>	8
1.2.6	6. Truthiness	8
1.2.7	7. Equality (<code>==</code>) vs Identity (<code>is</code>)	8
1.3	Examples	8
1.3.1	Level 1 — Basic (Type Checking)	8
1.3.2	Level 2 — Intermediate (Truthiness)	8
1.3.3	Level 3 — Hard (Equality vs Identity with None)	8
1.4	Common Mistakes	9
1.5	Exercises	9
2	Strings and Manipulation	11
2.1	Why This Matters	11
2.2	Explanation	11
2.2.1	1. Indexing and Slicing	11
2.2.2	2. Essential String Methods	11
2.2.3	3. F-Strings and The Format Mini-Language	12
2.3	Common Mistakes	12
2.4	Solutions to Exercises	12
3	Lists	15
3.1	Why This Matters	15
3.2	Explanation	15
3.2.1	1. Indexing, Slicing, and Mutating	15
3.2.2	2. Essential List Methods	15
3.2.3	3. Nested Lists (2D Data)	16
3.3	Common Mistakes: The Aliasing Trap	16
3.4	Solutions to Exercises	16
4	Tuples	19
4.1	Why This Matters	19
4.2	Explanation	19
4.2.1	1. Syntax and Core Methods	19
4.2.2	2. Unpacking	19

4.3	Common Mistakes	20
4.3.1	The Single Element Tuple Trap	20
4.3.2	Nested Mutability	20
4.4	Solutions to Exercises	20
5	Sets	21
5.1	Why This Matters	21
5.2	Explanation	21
5.2.1	1. Syntax and Core Methods	21
5.2.2	2. Mathematical Operations	21
5.2.3	3. The Hashable Requirement	22
5.3	Solutions to Exercises	22
6	Dictionaries	23
6.1	Why This Matters	23
6.2	Explanation	23
6.2.1	1. Syntax and Core Methods	23
6.2.2	2. Safe Lookups	24
6.2.3	3. Iterating through a Dictionary	24
6.2.4	4. Nested Dictionaries	24
6.3	Solutions to Exercises	24

Part I

Python Fundamentals

Chapter 1

Python Data Types

1.1 Why This Matters

In Machine Learning, your data will take many forms: numerical arrays, text labels, boolean masks, and configurations. Understanding the core Python data types is essential because applying the wrong operation to the wrong type (e.g., trying to add a string to an integer) will crash your pipeline. Knowing how to check types and understand "truthiness" will help you debug faster.

1.2 Explanation

1.2.1 1. Numeric Types

- **int**: Whole numbers. Python handles arbitrarily large integers automatically.
- **float**: Decimal numbers. Standard for ML model weights and probabilities.
- **complex**: Complex numbers (e.g., $3 + 4j$). Rarely used in standard ML, but important for signal processing.

1.2.2 2. Text and Binary Types

- **str**: Text data, enclosed in quotes. Immutable.
- **bytes** / **bytearray**: Raw byte data. Useful for reading images, audio, or serialized ML models.

1.2.3 3. Boolean and None

- **bool**: **True** or **False**. Used for logic and masking.
- **None**: A special constant representing the absence of a value (similar to **null** in other languages).

1.2.4 4. Collection Types (Brief intro)

- **list**: Ordered, mutable sequence `[1, 2, 3]`.
- **tuple**: Ordered, immutable sequence `(1, 2, 3)`.
- **set**: Unordered collection of unique items `{1, 2, 3}`.
- **dict**: Key-value pairs `{"a": 1, "b": 2}`.

1.2.5 5. `type()` and `isinstance()`

- `type(x)` returns the exact type of `x`.
- `isinstance(x, type)` checks if `x` is an instance of a specific type (or a tuple of types). This is the preferred, professional way to check types.

1.2.6 6. Truthiness

In Python, every object can be evaluated as a boolean.

- **Falsy:** `0`, `0.0`, `""` (empty string), `None`, `[]`, `{}`, `()`, `set()`.
- **Truthy:** Almost everything else!

1.2.7 7. Equality (`==`) vs Identity (`is`)

- `==` checks if the **values** are equal.
- `is` checks if they are the **exact same object in memory**.

1.3 Examples

1.3.1 Level 1 — Basic (Type Checking)

```

1 x = 3.14
2 print(type(x))                # Output: <class 'float'>
3 print(isinstance(x, float))    # Output: True
4 print(isinstance(x, (int, float))) # Output: True (checks if int OR float)

```

1.3.2 Level 2 — Intermediate (Truthiness)

```

1 empty_list = []
2 if empty_list:
3     print("This won't print because empty lists are falsy.")
4
5 if not empty_list:
6     print("This WILL print!") # Preferred way to check if a list is empty

```

1.3.3 Level 3 — Hard (Equality vs Identity with None)

```

1 a = None
2 b = None
3
4 # Because None is a singleton (only one exists in memory), BOTH are True.
5 print(a == b) # True
6 print(a is b) # True
7
8 list_1 = [1, 2]
9 list_2 = [1, 2]
10 print(list_1 == list_2) # True, they have the same values
11 print(list_1 is list_2) # False, they are different objects in memory!

```


1.4 Common Mistakes

WARNING:

1. Using `type(x) == list` instead of `isinstance(x, list)`. `isinstance` gracefully handles inheritance (subclasses), which is crucial when working with ML libraries that create custom object types.
2. Checking if a list is empty using `if len(my_list) == 0:`. The Pythonic way is simply `if not my_list:`.

ML CONNECTION:

When reading a CSV dataset using Pandas, missing values often start as `None` or `NaN` (Not a Number, which is a float). Knowing how to detect and handle these null types, and understanding how booleans act as masks (e.g., selecting all rows where `age < 30`), is the absolute foundation of data preprocessing.

1.5 Exercises

Exercise 1: Truthiness and Types Predict the output of the following code snippet:

```
1 data = ""
2 result = 100
3
4 if data:
5     result = result + 50
6 else:
7     result = result - 20
8
9 print(result)
```

Solution: The output is 80. An empty string evaluates to False in Python, triggering the else block.

Exercise 2: Debugging and Best Practices The following function tries to check if the input is a string and if it is not empty. However, it uses poor Python practices. Rewrite it using `isinstance()` and Pythonic truthiness.

```
1 def process_text(text):
2     if type(text) == str:
3         if len(text) > 0:
4             print("Processing:", text)
5         else:
6             print("Empty string provided.")
7     else:
8         print("Error: Input is not a string.")
```

Refactored Solution:

```
1 def process_text(text):
2     if isinstance(text, str):
3         if text:
4             print("Processing:", text)
5         else:
6             print("Empty string provided.")
7     else:
8         print("Error: Input is not a string.")
```


Chapter 2

Strings and Manipulation

2.1 Why This Matters

In Machine Learning and Data Science, a massive portion of data is unstructured text (e.g., medical records, tweets, reviews, log files). To perform Natural Language Processing (NLP) or basic data cleaning, you must know how to manipulate strings efficiently.

2.2 Explanation

A string (`str`) in Python is a sequence of characters used to represent text.

IMPORTANT:

Immutability: Strings are immutable, meaning once created, they cannot be changed in place. Any operation that modifies a string actually creates a brand new string.

2.2.1 1. Indexing and Slicing

Because strings are sequences, you can access individual characters using an index (starting at 0) and slice them to extract substrings using `[start:stop:step]`.

```
1 text = "Machine Learning"
2
3 print(text[0])    # 'M'
4 print(text[-1])   # 'g' (negative indices count from the end)
5
6 # Slicing [start:stop] (Stop is exclusive)
7 print(text[0:7])  # 'Machine'
8 print(text[:7])   # 'Machine'
9 print(text[8:])   # 'Learning'
```

2.2.2 2. Essential String Methods

- `.lower()` / `.upper()`: Changes case.
- `.strip()`: Removes leading and trailing whitespace.
- `.split(delimiter)`: Splits the string into a list.
- `.replace(old, new)`: Replaces substrings.

```

1 # Cleaning dirty data by chaining methods
2 raw_data = "    user@email.com    \n"
3 clean_data = raw_data.strip().lower()
4 print(clean_data)    # 'user@email.com'

```

2.2.3 3. F-Strings and The Format Mini-Language

F-strings (formatted string literals) are the modern way to inject variables into strings. When inside an f-string's curly braces {}, anything to the **right** of the colon : is the formatting rule.

- {epoch:02d} (Integer Padding): Pads the integer with zeros until it is at least 2 characters wide.
- {loss:.3f} (Float Rounding): Rounds a floating point number to exactly 3 decimal places.
- {accuracy:.1%} (Percentage Conversion): Multiplies the float by 100, rounds to 1 decimal place, and appends a % sign.

```

1 epoch = 5
2 loss = 0.03456
3 acc = 0.892
4
5 print(f"[Epoch {epoch:02d}] Loss: {loss:.3f} | Acc: {acc:.1%}")
6 # Output: [Epoch 05] Loss: 0.035 | Acc: 89.2%

```

2.3 Common Mistakes

WARNING:

Trying to change a string character directly like `text[0] = 'X'` will throw a `TypeError`. Also, forgetting that `.strip()` does not remove spaces *inside* the string, only at the edges.

ML CONNECTION:

When preparing text data for an ML model (like a text classifier or an LLM), you usually build a **preprocessing pipeline**. This involves converting all text to lowercase, removing punctuation via `.replace()`, and splitting sentences into individual words via `.split()`.

2.4 Solutions to Exercises

Exercise 1: String Slicing

```

1 file_path = "/data/images/dataset_2024.csv"
2 print(file_path[-16:])

```

Solution: Negative indexing ensures we get exactly the last 16 characters regardless of what the parent directories are named.

Exercise 2: Data Cleaning

```
1 raw_tweet = "    Python is AWESOME for Machine Learning!!! \n"  
2 clean_tweet = raw_tweet.strip().lower().replace("!!!", "")  
3 print(clean_tweet)
```

Solution: Chaining methods works perfectly here. python is awesome for machine learning

Exercise 3: F-Strings and Formatting

```
1 epoch = 5  
2 loss = 0.0345678  
3 val_accuracy = 0.892  
4 print(f"[Epoch {epoch:02d}] Training Loss: {loss:.3f} | Validation Acc:  
    {val_accuracy:.1%}")
```

Solution: [Epoch 05] Training Loss: 0.035 | Validation Acc: 89.2%

Chapter 3

Lists

3.1 Why This Matters

In Machine Learning, lists are the absolute foundation for storing sequences of data. Before you learn advanced numerical arrays (like NumPy), you will use standard Python lists to store model predictions, sequences of tokens in NLP, or rows of data from a CSV file.

3.2 Explanation

A list in Python is an ordered, **mutable** (changeable) collection of items. Lists can contain items of different data types, including other lists. They are created using square brackets `[]`.

Because they are ordered sequences, you can index and slice them exactly like Strings. But unlike Strings (which are immutable), you can modify lists *in place*.

3.2.1 1. Indexing, Slicing, and Mutating

```
1 fruits = ["apple", "banana", "cherry"]
2
3 # Indexing and Slicing (Same as strings!)
4 print(fruits[0])      # 'apple'
5 print(fruits[-1])     # 'cherry'
6 print(fruits[:2])     # ['apple', 'banana']
7
8 # Mutating (Changing items in place)
9 fruits[1] = "blueberry"
10 print(fruits)        # ['apple', 'blueberry', 'cherry']
```

3.2.2 2. Essential List Methods

- `.append(item)`: Adds an item to the **end** of the list.
- `.insert(index, item)`: Inserts an item at a specific position.
- `.pop(index)`: Removes and returns the item at the index (defaults to the last item if no index is given).
- `.remove(value)`: Removes the *first* occurrence of a specific value.
- `.extend(list2)`: Adds all elements of a second list to the end of the first.
- `.sort()`: Sorts the list in place (modifies the original list).

3.2.3 3. Nested Lists (2D Data)

Lists can contain other lists. This is how we represent 2D data (like a spreadsheet or an image) in pure Python.

```

1 matrix = [
2     [1, 2, 3],
3     [4, 5, 6],
4     [7, 8, 9]
5 ]
6
7 # To get the number 6:
8 # First get the middle row (index 1), then the last item (index -1)
9 print(matrix[1][-1]) # 6

```

3.3 Common Mistakes: The Aliasing Trap

WARNING:

This is the most common mistake made by Python beginners. Because lists are mutable, variables simply **point** to the list in memory.

If you write `b = a`, you did not copy the list. You created a second variable pointing to the same list. Modifying `b` will modify `a`.

To make an actual independent copy, use the `.copy()` method or slice it `[:]`:

```

1 a = [1, 2, 3]
2 c = a.copy()
3 c.append(5)
4 print(a)          # Still [1, 2, 3], 'c' is independent.

```

ML CONNECTION:

When building neural networks from scratch, the architecture is often defined as a list of layers. Furthermore, when reading raw data before passing it to a library like Pandas, it usually starts as a "list of lists" where each inner list represents a single row in a dataset.

3.4 Solutions to Exercises

Exercise 1: The Queue

```

1 queue = ["img1.png", "img2.png"]
2 queue.append("img3.png")
3 print(queue.pop(0))

```

Solution: `.append("img3.png")` safely adds the new image to the very end of the list. `.pop(0)` removes the element at index 0 (the first element) and returns it.

Exercise 2: Nested Data Extraction

```

1 batch = [
2     [ [0, 0], [0, 0] ],
3     [ [1, 1], [1, 1] ],
4     [ [2, 2], [2, 9] ]
5 ]
6 print(batch[2][1][1])

```


Solution: `batch[2]` selects the 3rd image. `[1]` selects the 2nd row `[2, 9]`. `[1]` selects the 2nd element 9.

Exercise 3: The Copy Trap

```
1 original_data = [100, 200, 300]
2 normalized = original_data.copy()
3 normalized.remove(300)
4
5 print("Original:", original_data)
6 print("Normalized:", normalized)
```

Solution: By using `.copy()`, Python creates a brand new list in memory. Now `normalized` points to this new list, while `original_data` points to the old one.

Chapter 4

Tuples

4.1 Why This Matters

If a tuple is just a list you can't change, why use it?

1. **Safety:** In Machine Learning, you often have data that should *never* change during execution (like the shape of an image matrix: (256, 256, 3)). Using a tuple guarantees that a stray line of code won't accidentally mutate this critical data.
2. **Speed and Memory:** Because tuples are immutable, Python optimizes them. They take up slightly less memory and are faster to iterate through than lists.
3. **Dictionary Keys:** Lists cannot be used as keys in a dictionary (because they can change), but tuples can.

4.2 Explanation

A tuple in Python is an ordered, **immutable** (unchangeable) collection of items. Tuples are created using parentheses () (though the parentheses are actually optional; it's the commas , that define a tuple). Because they are ordered, you index and slice them just like lists and strings.

4.2.1 1. Syntax and Core Methods

```
1 dimensions = (1920, 1080)
2 hyperparameters = 0.01, 32, 64 # Parentheses are optional!
3
4 print(dimensions[0])           # 1920
5 # dimensions[0] = 800 <--- TypeError!
```

Because they are immutable, tuples only have **two** methods: `.count(value)` and `.index(value)`.

4.2.2 2. Unpacking

A very Pythonic feature: you can assign multiple variables at once from a tuple.

```
1 model_config = ("Adam", 0.001, 100)
2 optimizer, learning_rate, epochs = model_config
3
4 print(learning_rate) # 0.001
```

4.3 Common Mistakes

4.3.1 The Single Element Tuple Trap

If you want to create a tuple with exactly ONE element, you **must** include a trailing comma. Parentheses alone are just treated as mathematical grouping!

```
1 math_variable = (5)
2 print(type(math_variable)) # <class 'int'>
3
4 single_tuple = (5,)
5 print(type(single_tuple)) # <class 'tuple'>
```

4.3.2 Nested Mutability

A tuple itself cannot be changed. However, if a tuple *contains* a mutable object (like a list), you CAN mutate the list inside the tuple!

```
1 complex_tuple = (1, 2, [3, 4])
2 complex_tuple[2].append(5)
3 print(complex_tuple) # (1, 2, [3, 4, 5])
```

Why? Because the tuple is still pointing to the EXACT same list object in memory. The tuple didn't change its pointers; the list itself changed.

ML CONNECTION:

When you use a library like NumPy or TensorFlow, the shape of a multi-dimensional array or tensor is ALWAYS returned as a tuple (e.g., (batch_size, height, width, channels)). Furthermore, many ML functions return multiple values (e.g., return loss, accuracy). When a function returns multiple values separated by commas, it is secretly returning a tuple!

4.4 Solutions to Exercises

Exercise 1: Unpacking

```
1 results = ("ResNet50", 0.92, 14.5)
2 name, acc, time = results
```

Solution: This unpacks the tuple into three separate variables seamlessly.

Exercise 2: The Tuple Trap

```
1 lr_tuple = (0.01,)
```

Solution: Without the comma, Python evaluates (0.01) as a standard float.

Exercise 3: Swapping Variables

```
1 x = 10
2 y = 20
3 x, y = y, x
```

Solution: This is the famous Pythonic variable swap. The right side packs a tuple (20, 10), and the left side immediately unpacks it, bypassing the need for a temporary variable.

Chapter 5

Sets

5.1 Why This Matters

Sets are incredibly powerful for two main reasons:

1. **Deduplication:** The fastest way to remove duplicates from a list is to convert it into a set.
2. **Membership Testing:** Because sets use a data structure called a "hash table" under the hood, checking if an item exists in a set (e.g., if "apple" in my_set:) is almost instantly fast, even if the set has millions of items. Doing this with a list is incredibly slow.

5.2 Explanation

A set in Python is an unordered, mutable collection of **unique** items. You can think of it exactly like a mathematical set. Because they are unordered, **they do not support indexing or slicing** (you cannot do my_set[0]).

5.2.1 1. Syntax and Core Methods

```
1 unique_numbers = {1, 2, 3, 4, 4, 4, 5}
2 print(unique_numbers) # {1, 2, 3, 4, 5} -> Duplicates removed!
3
4 # Creating an empty set
5 empty_set = set() # MUST use set(), because {} creates a dict!
```

- .add(item): Adds an item to the set.
- .remove(item): Removes an item (throws an error if it doesn't exist).
- .discard(item): Removes an item (does NOTHING if it doesn't exist - safer!).

5.2.2 2. Mathematical Operations

Sets shine when you need to compare two groups of data.

```
1 group_a = {"apple", "banana", "cherry"}
2 group_b = {"cherry", "date", "elderberry"}
3
4 print(group_a | group_b) # Union: {'apple', 'banana', 'cherry', 'date', 'elderberry'}
```

```

5 print(group_a & group_b) # Intersection: {'cherry'}
6 print(group_a - group_b) # Difference: {'apple', 'banana'}

```

5.2.3 3. The Hashable Requirement

Sets can only store **immutable** (hashable) items. This means you can put strings, integers, and tuples inside a set, but you CANNOT put a list or a dictionary inside a set!

```

1 valid_set = {1, "hello", (10, 20)} # Works perfectly
2 # invalid_set = {1, "hello", [10, 20]} <--- TypeError!

```

ML CONNECTION:

In Natural Language Processing (NLP), you often need to find the “vocabulary” of a document—the unique words used. If a document has 10,000 words, you simply run `vocab = set(words_list)` to instantly get the unique words. Sets are also used to quickly identify overlapping features between datasets.

5.3 Solutions to Exercises

Exercise 1: Deduplication

```

1 emails = ["a@a.com", "b@b.com", "a@a.com", "c@c.com", "b@b.com"]
2 unique_emails = set(emails)
3 print(len(unique_emails)) # 3

```

Exercise 2: Intersection

```

1 user_likes = {"The Matrix", "Inception", "Interstellar"}
2 friend_likes = {"Interstellar", "The Prestige", "The Matrix"}
3 print(user_likes & friend_likes)

```

Exercise 3: The Empty Set Trap

```

1 unique_ips = set()
2 unique_ips.add("192.168.1.1")

```

Solution: `{}` evaluates to an empty dictionary, which does not have an `.add()` method. You must use `set()`.

Chapter 6

Dictionaries

6.1 Why This Matters

In Machine Learning, dictionaries are the absolute backbone of data structure formatting.

- **JSON Data:** APIs and datasets often return data in JSON format, which maps perfectly 1-to-1 with Python dictionaries.
- **Hyperparameters:** When training an ML model, you will often pass in a dictionary of configuration settings (e.g., `{"learning_rate": 0.01, "batch_size": 32}`).
- **Model Weights:** PyTorch and TensorFlow store neural network layers and weights inside dictionaries.

6.2 Explanation

A dictionary (dict) in Python is an unordered, **mutable** collection of **key-value pairs**. It is incredibly fast for looking up values based on a unique key.

Keys MUST be unique and immutable (like strings, integers, or tuples). **Values can be absolutely anything.**

6.2.1 1. Syntax and Core Methods

```
1 student = {  
2     "name": "Alice",  
3     "age": 28,  
4     "courses": ["Math", "CompSci"]  
5 }  
6  
7 # Modifying or Adding a new key-value pair  
8 student["age"] = 29
```

- `.get(key, default)`: Safely gets a value without crashing if it doesn't exist.
- `.keys()`: Returns all keys.
- `.values()`: Returns all values.
- `.items()`: Returns all (key, value) tuples.
- `.update(dict2)`: Merges another dictionary into this one.

6.2.2 2. Safe Lookups

Accessing a key that doesn't exist using brackets `[]` will crash your program with a `KeyError`. Always use `.get()` if you aren't 100% sure the key exists.

```
1 config = {"theme": "dark"}
2 # print(config["font_size"]) <--- KeyError! Program crashes.
3
4 font = config.get("font_size", 12) # Safe
```

6.2.3 3. Iterating through a Dictionary

To loop through a dictionary effectively, use the `.items()` method, which unpacks the key and value simultaneously.

```
1 model_metrics = {"accuracy": 0.95, "loss": 0.02}
2
3 for metric_name, value in model_metrics.items():
4     print(f"{metric_name}: {value}")
```

6.2.4 4. Nested Dictionaries

When working with JSON data, you will often deal with dictionaries inside dictionaries.

```
1 dataset = {
2     "user_101": {
3         "name": "Bob",
4         "purchases": [15.99, 20.00]
5     }
6 }
7 print(dataset["user_101"]["purchases"][0]) # 15.99
```

ML CONNECTION:

When you use Pandas, creating a `DataFrame` (a table of data) from scratch is often done using a dictionary where the keys are the column names and the values are lists representing the rows.

6.3 Solutions to Exercises

Exercise 1: Safe Access

```
1 user_profile = {"name": "Sarah", "email": "sarah@example.com"}
2 phone = user_profile.get("phone", "No phone provided")
```

Solution: `.get()` checks if the key "phone" exists. It safely returns the fallback string instead of crashing.

Exercise 2: Dictionary Updating

```
1 default_config = {"lr": 0.01, "batch_size": 32, "epochs": 100}
2 user_config = {"batch_size": 64, "optimizer": "Adam"}
3 default_config.update(user_config)
```

Solution: The `.update()` method merges `user_config` into `default_config`, overwriting existing keys and adding new ones.

Exercise 3: Nested Extraction


```
1 print(api_response["data"]["forecast"][1]["conditions"])
```

Solution: Chained indexing goes from outer dictionary ["data"] to inner list ["forecast"], to index [1], to key ["conditions"].

Part II

Core Python

Chapter 7

Control Flow

7.1 Why This Matters

Control flow is how you dictate the "path" your code takes. Until now, our code has executed strictly from top to bottom, line by line. Control flow allows your code to make **decisions** (**if**), **repeat** tasks (**for**, **while**), and skip lines of code entirely based on logic.

In Machine Learning, control flow is used everywhere:

- **if statements:** "If the model's accuracy drops below 80%, stop training."
- **for loops:** "For every image in the dataset, resize it to 256x256."
- **while loops:** "While the loss function is still decreasing, keep training."

7.2 The if Statement

Python uses **if**, **elif** (else if), and **else**. It relies entirely on **indentation** to know what code belongs inside the block.

```
1 accuracy = 0.85
2
3 if accuracy > 0.90:
4     print("Deploy the model!")
5 elif accuracy > 0.80:
6     print("Keep training.")
7 else:
8     print("Failing.")
```

*Note: You can chain logical operators like **and**, **or**, and **not**.*

7.3 The for Loop

A **for** loop in Python iterates directly over the items in a collection (like a list, tuple, string, or dictionary).

```
1 # Iterating over a list
2 loss_values = [0.5, 0.4, 0.2]
3 for loss in loss_values:
4     print(f"Current loss: {loss}")
5
6 # Generating number sequences
7 for epoch in range(5):
8     print(f"Training epoch {epoch}")
```

7.4 The while Loop

A **while** loop runs forever as long as its condition evaluates to **True**. You **MUST** ensure the condition eventually becomes **False**, or you will create an **infinite loop**.

```
1 battery = 100
2 while battery > 0:
3     battery -= 20
```

7.5 Loop Control

- **break**: Completely shatters the loop. The loop ends immediately.
- **continue**: Skips the rest of the *current* iteration and instantly jumps to the *next* iteration.

ML CONNECTION:

When you build neural networks using PyTorch from scratch, the core of your program is quite literally just a massive **for** loop (iterating over epochs), nested inside another **for** loop (iterating over batches of data)!

7.6 Solutions to Exercises

Exercise 1: The Classifier

```
1 if age < 18:
2     print("Minor")
3 elif age >= 18 and age <= 64:
4     print("Adult")
5 else:
6     print("Senior")
```

Exercise 2: The Data Cleaner

```
1 readings = [12.5, 14.1, None, 11.2, None, 15.0]
2 for result in readings:
3     if result is None:
4         continue
5     print(result)
```

Exercise 3: Early Stopping

```
1 loss = 1.0
2 epoch = 0
3 while True:
4     epoch += 1
5     loss -= 0.1
6     if loss < 0.3 or epoch == 5:
7         break
```